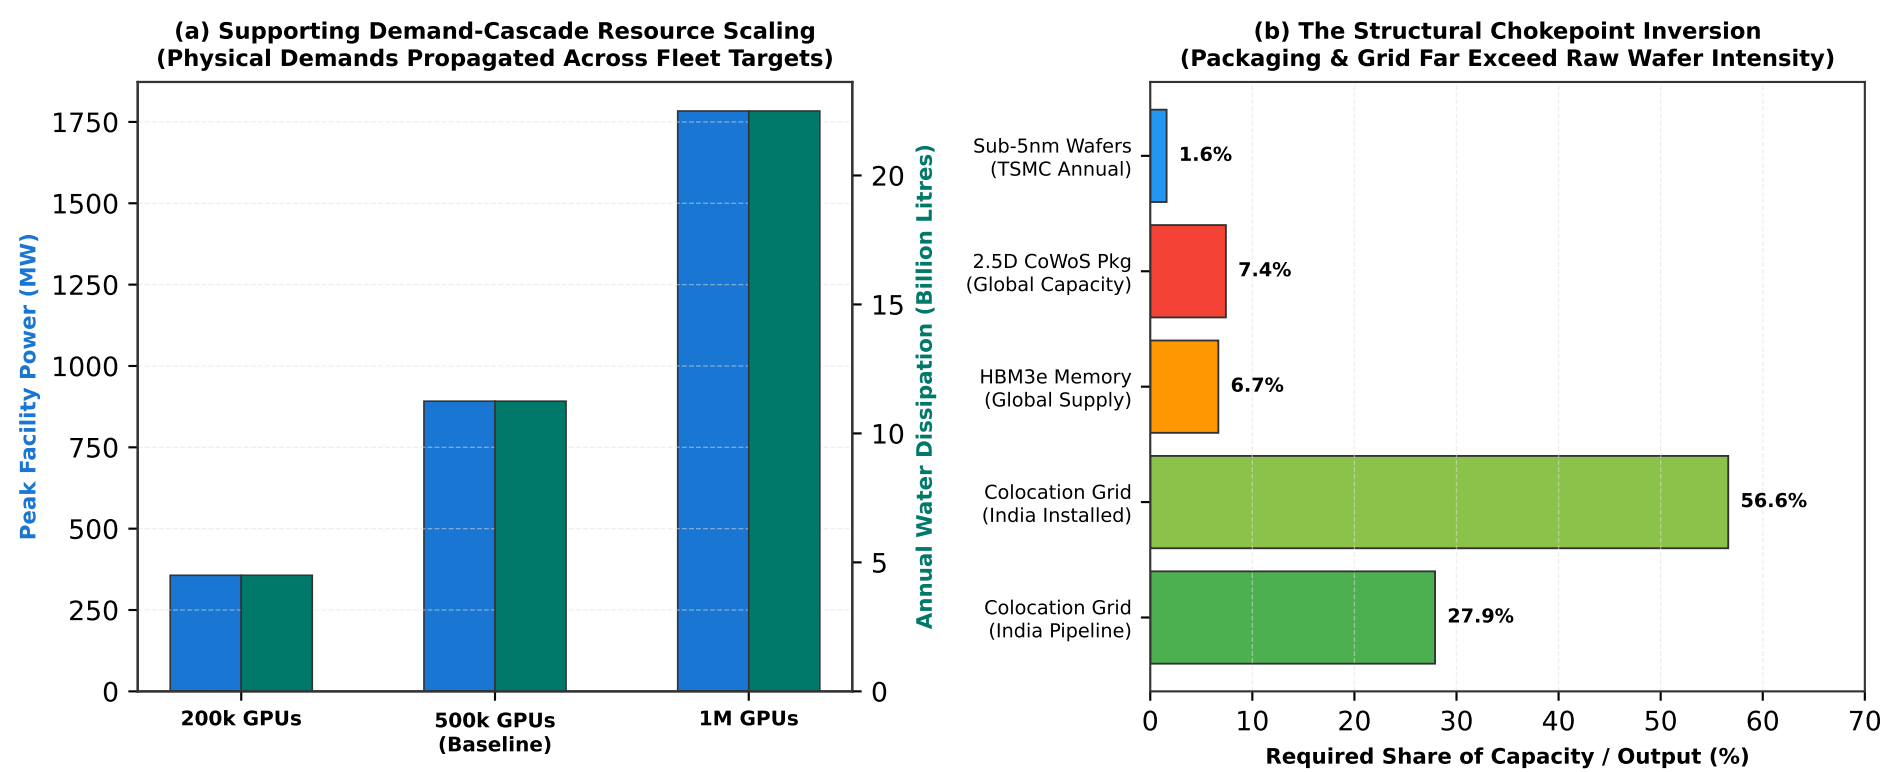
(a) Supporting Demand-Cascade Resource Scaling
(Physical Demands Propagated Across Fleet Targets)
Peak Facility Power (MW)
Annual Water Dissipation (Billion Litres)
0
250
500
750
1000
1250
1500
1750
0
5
10
15
20
200k GPUs
500k GPUs
(Baseline)
1M GPUs
(b) The Structural Chokepoint Inversion
(Packaging & Grid Far Exceed Raw Wafer Intensity)
Sub-5nm Wafers
(TSMC Annual)
1.6%
2.5D CoWoS Pkg
(Global Capacity)
7.4%
HBM3e Memory
(Global Supply)
6.7%
Colocation Grid
(India Installed)
56.6%
Colocation Grid
(India Pipeline)
27.9%
0
10
20
30
40
50
60
70
Required Share of Capacity / Output (%)